

# Erlang *Dos* and *Don'ts*

Richard Carlsson

Klarna

# Prefer to avoid macros

- **Don't:**

- define(msg(X), {message, self(), X}).

- **Do:**

- msg(X) → {message, self(), X}.

- **Why:**

- Macros shared between modules require recompiling of all modules if you make a change
  - You get better compiler warnings with functions
  - The compiler can handle inlining for you

# Macros should be well formed

- **Don't:**

- define(p(X), io:write(X), ).
  - define(do(X), fun () -> ?p(X) ).
  - define(done(X), ?p(X) ok end ).

- **Do:**

- % Something else - anything else

- **Why:**

- Screws up readability of the code
  - Makes it impossible to parse in EDoc and other tools that try to analyze the source code
  - There is a circle in hell reserved for those who use macros this way

# Do not export everything

- **Don't:**

- `compile(export_all)`.

- **Do:**

- `export([...])`.

- **Why:**

- Clear distinction between official API and internal support functions
  - Lets the compiler make better optimizations
  - Helps Dialyzer make better type analysis

# List all cases explicitly

- **Don't:**

```
case is_foo(X) of
  true -> ... ;
  _     -> ...
end
```

- **Do:**

```
case is_foo(x) of
  true   -> ... ;
  false  -> ...
end
```

- **Why:**

- If there's a bug in `is_foo(X)` that makes it return a non-boolean value sometimes, you will not notice it, because the catch-all case will always be selected

# Be careful with list\_to\_atom()

## ■ Don't:

```
foo(InputString) ->  
    ... list_to_atom(InputString)
```

## ■ Do:

```
foo("on") -> on;  
foo("off") -> off;  
...
```

## ■ Why:

- The atom table in Erlang is limited in size, and is never garbage collected
- A malicious client could send random strings until the table is full, crashing our node
- list\_to\_existing\_atom(String) can also be used

# Don't create improper lists

- **Don't:**

`L1 = [L | 42]    % 42 as "tail"`

- **Do:**

`L1 = L ++ [42]    % if you have to`

`L1 = [L, 42]    % if a deep list is ok`

- **Why:**

- Almost all functions that expect a list will crash if the list is not proper (i.e., does not end with `[]`)
- Even if you don't care right now, someone may want to use a normal list function on `L1` later
- It will make type analysis in Dialyzer less exact

# Avoid length(List)

- **Don't:**

```
foo(L) when length(L) > 0 ->  
...
```

- **Do:**

```
foo([_|_]=L) ->  
...
```

- **Why:**

- length(L) has to traverse the entire list L
- If you do this within a loop over L, it will cause quadratic time complexity
- Pattern matching instead takes constant time



# Don't append to the right

- **Don't:**

```
foo([Ls | Lss], ..., As) ->  
    ... foo(Lss, ..., As ++ Ls);  
foo([], As) ->  
    As.
```

- **Do:**

```
foo([Ls|Lss], ..., As) ->  
    ... foo(Lss, ..., reverse(Ls) ++ As);  
foo([], As) ->  
    reverse(As).
```

- **Why:**

- A ++ B creates a new list by copying A onto B, which takes time proportional to the length of A
- If you do this within a loop over A, it will cause quadratic time complexity

# Avoid --

- **Don't:**

```
Rest = BigList -- AnotherBigList
```

- **Do:**

```
S1 = ordsets:from_list(List1),  
S2 = ordsets:from_list(List2),  
Rest = ordsets:subtract(S1, S2)
```

- **Why:**

- A -- B searches and rewrites A once for each element in B, so it runs in  $O(A*B)$  time
  - If B is guaranteed to be short, this is OK
- $[1,2,2] -- [2] \rightarrow [1,2]$ 
  - If it's necessary to preserve order and repeated elements in A, then -- can be OK

# Updating tuples means copying

- **Don't:**

```
T1 = setelement(999, T0, Value)
```

- **Do:**

```
A0 = array:new(),  
A1 = array:set(999, Value, A0)
```

- **Why:**

- Updating a tuple of size N copies N machine words of data, just to replace a single word
- The array module is better if N is large (>50), and it also handles sparse arrays well
- Remember that records are tuples

# Use the right size operator

- **Don't:**

```
foo(T, B) when size(T) > 2 ->  
    Bytes = size(B)
```

- **Do:**

```
foo(T, B) when tuple_size(T) > 2 ->  
    Bytes = byte_size(B),  
    Bits = bit_size(B)
```

- **Why:**

- size(X) is too overloaded (accepts both tuples and binaries), giving no clues to the reader, nor to the compiler or Dialyzer
- Note that if e.g. B = <<255, 255, 17:3>>, then bit\_size(B) is 19, and byte\_size(B) is 3 – the smallest number of bytes needed

# Avoid is\_record()

- **Don't:**

```
foo(R) when is_record(R, bar) ->  
... X = R#bar.foo
```

- **Do:**

```
foo(#bar{foo=X}=R) ->  
...
```

- **Why:**

- More obvious to the reader
- Easier to find uses of #bar with grep etc.
- Better hints to the compiler

# Store text as binaries

- **Don't:**

% huge dict of strings in server state  
D1 = dict:store(Key, String, D0)

- **Do:**

% encode as UTF-8 (accepts deep lists)  
B = unicode:characters\_to\_binary(Txt)

- **Why:**

- A string (list of characters) uses 8 bytes per character in RAM (16 bytes on 64-bit machines)
- UTF-8 binary strings are much more compact
- Use strings as temporary buffers and use binaries for more long-term storage
- BUT: not worth the effort if it's not a lot of data

# Use I/O-lists for output

- **Don't:**

```
Out = L1 ++ L2,  
io:format("Result: ~s\n", [Out])
```

- **Do:**

```
Out = [L1, L2],  
io:format("Result: ~s\n", [Out])
```

- **Why:**

- All functions that send data to output streams accept I/O-lists as data
- It's a complete waste of time to create a flat list if you don't really need it to be flat
- No need to turn binaries to lists first – you can mix binaries and character codes in I/O-lists

# Parameter type overload

- **Don't:**

```
foo(X) when is_integer(X) ->  
    -X;  
foo(X) when is_float(X) ->  
    1 / X.
```

- **Do:**

```
negate(X) -> -X.
```

```
inverse(X) -> 1 / X.
```

- **Why:**

- The reader doesn't have to know the type of X to know what the function does
- Allow compiler to make better optimizations and Dialyzer to make better type analysis



# Confusing types

- **Don't:**

```
% integer for age or string for name  
-spec f() -> integer() | string().
```

- **Do:**

```
-spec f() -> {age, integer()}  
          | {name, string()}.
```

- **Why:**

- Tagging makes representation self-explaining
- Much easier to debug
- Better type analysis in Dialyzer

# Avoid leaking data representation

- **Don't:**

- spec f() -> {user, atom(), string()}.

- **Do:**

- spec f() -> {user, [Info]}  
when Info :: {name, string()}  
          | {id, atom()}

- **Why:**

- Makes it hard to change representation later
  - Is all the code that knows about the representation under your control?
  - How many lines of code need changing, and in how many different modules?
  - What about data stored on disk, or in an ETS table, or in the state of a server process?

# Spawn and link

- **Don't:**

`spawn(worker, start, [Data])`

- **Do:**

`spawn_link(worker, start, [Data])`

- **Why:**

- Ensures all linked processes get cleaned up, if one of them crashes
- Orphaned processes can be hard to track down
- Only spawn without linking if you *know* that is exactly what you want to do
- Try to ensure that unlinked processes are registered somehow, so you can find them later

# Make code upgrade possible

- **Don't:**

```
loop(...) ->  
    receive  
    ... -> ... loop(...)  
end.
```

- **Do:**

```
loop(...) ->  
    receive  
    ... -> ... ?MODULE:loop(...)  
end.
```

- **Why:**

- If you only loop using local calls, the process will not switch to the new code when you reload the module. Reloading twice will kill the process.

# Anonymous funs should be short-lived

- **Don't:**

```
F = fun (X) -> X + N end,  
ets:insert(Table, {Key, F})
```

- **Do:**


```
F = fun mymod:myfun/1,  
ets:insert(Table, {Key, F, N})
```

- **Why:**

- If the module that the anonymous fun belongs to gets reloaded more than once, the existing fun-values will become invalid and will cause the processes that hold them to be killed
- Named funs do not have this effect

# OTP

**Read it!**



# Use the application environment

- **Don't:**

```
init([]) ->  
    {ok, #state{port = 1234}}.
```

- **Do:**

```
init([]) ->  
    Port = application:get_env(port),  
    {ok, #state{port = Port}}.
```

- **Why:**

- Makes it easy to reconfigure your application via the system config file.
- Inject dependencies between applications via configuration, instead of hard-coding them
- Preferably, use `get_env` at startup time only, and use messages for dynamic reconfiguration

# Use gen\_server

- **Don't:**

```
% home made server loop
start(S) -> spawn(?MODULE, loop, [S]).
loop(S0) ->
    receive
        ...
        loop(S1)
    end.
```

- **Do:**

```
start(...) -> gen_server:start_link(...).
```

- **Why:**

- Predictability, easier to understand and debug.
- Don't repeat the same old mistakes as everybody else (at least not on company time)



# Use supervisors dynamically

- **Don't:**

- % non-supervisor spawning a worker  
spawn\_link(worker, run, [Data])

- **Do:**

- % use a 'simple\_one\_for\_one' supervisor  
% as a process factory with supervision

- supervisor:start\_child(worker\_sup, [Data])

- **Why:**

- Easier to control, debug, and visualize in tools
  - Better crash reports
  - Automatic restarts

# Use proc\_lib to spawn

- **Don't:**

```
spawn_link(worker, run, [Data])
```

- **Do:**

```
% if you really want to use spawn  
% (mainly for simple background jobs)
```

```
proc_lib:spawn_link(worker, run, [Data])
```

- **Why:**

- Sets up OTP environment in the new process (certain process dictionary settings used by OTP)
- Better crash reports via SASL
- Easier for OTP tools to visualize etc.

# Keep logic out of supervisors

- **Don't:**

```
init([]) ->  
    application:start(crypto),  
    Children = case ... of  
    ...
```

- **Do:**

```
init([]) ->  
    Children = [...],  
    {ok, {{one_for_one,3,10}, Children}}.
```

- **Why:**

- If something stops working, your whole application might fail to start, or cannot restart
- This kills the node
- Preferably, don't even call other modules

# The init() callback must finish fast

- **Don't:**

```
init([LSocket]) ->  
    {ok, S} = gen_tcp:accept(LSocket),  
    {ok, #state{socket=S}}.
```

- **Do:**

```
% this uses the server timeout trick  
init([A, B]) ->  
    {ok, #state{a=A, b=B}, 0}.
```

```
handle_info(timeout, State) -> ...
```

- **Why:**

- The process trying to start your gen\_server can time out and crash if init() takes too long time

# Make handle\_...() return quickly

- **Don't:**

```
handle_call({find, X}, _From, St) ->
    {reply, {ok, search_db(X)}, St}.
```

- **Do:**

```
handle_call({find, X}, From, St) ->
    proc_lib:spawn_link(fun ()->
        Reply = {ok, search_db(X)},
        gen_server:reply(From, Reply))
    end),
    {noreply, St}.
```

- **Why:**

- Blocks the gen\_server process from handling other requests until the computation is done

# Maybe make it asynchronous

- **Don't:**

```
handle_call({prompt, P}, _From, St) ->  
    {reply, {ok, user_input(P)}, St}
```

- **Do:**

```
handle_cast({prompt, P, From}, St) ->  
    proc_lib:spawn_link(fun ()->  
        From ! {prompt, user_input(P)},  
        end),  
    {noreply, St}.
```

- **Why:**

- The caller is blocked, and could time out and crash if the call takes a long time.
- Cast + Reply is more complicated for the caller, but lets it do other things while waiting.

Here endeth the lesson